

Boolean Algebra over Trapdoor Sets: A Practical Framework for Privacy-Preserving Set Operations with Probabilistic Guarantees

Alexander Towell
Department of Computer Science
Southern Illinois University Edwardsville
Email: atowell@siue.edu

Abstract—We present Hash-Based Trapdoor Sets (HBTS), a practical framework for privacy-preserving set operations that combines cryptographic hash functions with probabilistic data structures. Unlike traditional approaches using fully homomorphic encryption or secure multi-party computation, HBTS achieves microsecond-scale performance by embracing approximate operations with explicitly managed error rates.

HBTS provides *trapdoor representation*¹ through one-way hash transformations—the underlying data cannot be recovered from hash values. Operations return *Bernoulli Booleans*: plaintext approximate results with explicit error rates (false positive rate α , false negative rate β). This layered model accepts observable query results in exchange for practical efficiency, making it suitable for applications where query patterns are not sensitive or results are shared.

Our framework provides: (1) a systematic approach to propagating error bounds through composed set operations using Bernoulli type composition rules, (2) efficient Boolean and symmetric difference operations on hash-transformed data with quantifiable false positive rates, and (3) practical implementations of privacy-preserving primitives including private set intersection and secure aggregation. We build upon established techniques from probabilistic data structures (Bloom filters, HyperLogLog) while adding cryptographic privacy through one-way hash transformations.

Since HBTS operations reduce to standard hash computations and set operations, they are orders of magnitude faster than homomorphic encryption or MPC approaches. The framework provides false positive rates bounded by hash collision probability (e.g., 2^{-256} for SHA-256), with privacy guarantees that depend on input entropy and threat model assumptions. We provide a reference Python implementation and validate HBTS through applications to private set intersection, secure deduplication, and federated learning aggregation, showing practical applicability where approximate results with explicit error bounds are acceptable.

Index Terms—cryptographic hash functions, trapdoor data structures, privacy-preserving computation, approximate algorithms, probabilistic data structures

I. INTRODUCTION

The proliferation of cloud computing and data analytics has created an urgent need for privacy-preserving computational frameworks that enable operations on sensitive data without exposing the underlying values. Traditional approaches to this problem fall into two categories: cryptographic techniques such as fully homomorphic encryption (FHE) [1] and secure multi-party computation (MPC) [2], and statistical techniques such as differential privacy [3]. While powerful, these approaches often suffer from computational overhead that limits their practical deployment.

We present Hash-Based Trapdoor Sets (HBTS), a practical framework that achieves efficient privacy-preserving set operations by combining cryptographic hash functions with probabilistic data structures. Our approach differs from existing solutions by explicitly embracing approximation, making error rates transparent throughout the computational pipeline. This design choice enables HBTS to achieve microsecond-scale performance while providing privacy bounded by hash collision probabilities.

The core insight is that many applications can tolerate controlled error rates in exchange for efficiency and privacy. Cryptographic hash functions serve as one-way transformations: sensitive data is mapped into a hash domain where equality testing is preserved probabilistically while original values remain computationally hidden. All subsequent operations work exclusively on hash values, never requiring access to plaintext.

HBTS embodies a *layered privacy model*: the hash representation is a *trapdoor* (underlying values cannot be recovered), while operations return *Bernoulli Booleans*—plaintext approximate results with explicit error rates. This separation is key to understanding HBTS’s guarantees: we hide what the data *is*, but not the pattern of query *results*. The framework draws on Bernoulli type theory [4], which provides systematic error propagation rules for composed approximate operations.

¹We use “trapdoor” to mean a one-way encoding keyed by a secret, not a trapdoor permutation in the standard cryptographic sense. The hash key k enables creation of encodings but does not permit recovery of inputs—there is no efficient inversion.

A. Motivating Example

Consider a healthcare consortium where multiple hospitals need to identify patients enrolled in overlapping clinical trials without revealing their complete patient lists. Traditional approaches would require either: (1) a trusted third party to perform the intersection, violating privacy requirements, (2) homomorphic encryption with prohibitive computational costs, or (3) complex MPC protocols requiring multiple rounds of communication.

Using HBTS, each hospital can:

- 1) Transform patient identifiers using a shared cryptographic hash function
- 2) Perform set intersection directly on hash values
- 3) Obtain results with explicit error bounds (e.g., hash collision rate $< 2^{-256}$)
- 4) Complete the entire operation in milliseconds rather than minutes

B. Contributions

This paper makes the following contributions:

- **Systematic Error Propagation:** We provide a formal framework for propagating error bounds through composed set operations on hash-transformed data, enabling applications to reason about accuracy-privacy trade-offs.
- **Reference Implementation:** We provide a Python reference implementation demonstrating that combining cryptographic hashing with probabilistic data structures yields practical privacy-preserving set operations using only standard-library primitives.
- **Integration of Existing Techniques:** We show how established algorithms (Bloom filters, HyperLogLog) can be enhanced with cryptographic privacy guarantees through systematic application of hash transformations.
- **Real-World Applications:** We validate HBTS through implementations of private set intersection, secure deduplication, and federated learning aggregation, demonstrating practical utility where approximate results are acceptable.
- **Security Analysis:** We provide formal analysis showing that privacy is bounded by the collision probability of the underlying hash function, with explicit quantification of error rates.

II. BACKGROUND AND THREAT MODEL

A. Cryptographic Hash Functions

A cryptographic hash function $H : \{0,1\}^* \rightarrow \{0,1\}^n$ maps arbitrary-length inputs to fixed-size outputs satisfying the standard properties: preimage resistance ($O(2^n)$ to invert), second preimage resistance ($O(2^n)$ to find a collision for a given input), and collision resistance ($O(2^{n/2})$ to find any collision). HBTS exploits these properties to create one-way transformations that preserve equality testing probabilistically while preventing recovery of original values.

B. Approximate Data Structures

Approximate data structures trade perfect accuracy for improved space or time complexity. The canonical example is the Bloom filter [5], which supports membership queries with false positives but no false negatives. HBTS builds upon these well-established techniques, adding cryptographic privacy through hash transformations while maintaining explicit error bounds.

Definition 1 (Bernoulli Boolean). A Bernoulli Boolean $\mathcal{B}(\mathbb{B})$ is an observed boolean $\tilde{b}$ with error rates (α, β) where:

- $\alpha = \mathbb{P}[\tilde{b} = \text{true} \mid b = \text{false}]$ (false positive rate)
- $\beta = \mathbb{P}[\tilde{b} = \text{false} \mid b = \text{true}]$ (false negative rate)

We distinguish second-order Bernoulli Booleans where $\alpha \neq \beta$ (asymmetric errors, as in Bloom filters and hash-based membership) from first-order where $\alpha = \beta$ (symmetric errors). HBTS produces second-order Bernoulli Booleans with $\alpha = 2^{-n}$ (hash collision) and $\beta = 0$ (no false negatives).

Definition 2 (Confusion Matrix). The error behavior of a Bernoulli Boolean is fully characterized by its confusion matrix:

$$Q = \begin{pmatrix} 1 - \alpha & \alpha \\ \beta & 1 - \beta \end{pmatrix}$$

where $Q_{ij} = \mathbb{P}[\tilde{b} = j \mid b = i]$ gives the probability of observing j given latent value i . Key properties:

- **Identity matrix** ($\alpha = \beta = 0$): perfect observation, no privacy
- **Rank-1 matrix** ($\alpha + \beta = 1$): complete information loss, optimal privacy
- **Rank-2 matrix** ($\alpha + \beta < 1$): information preserved, privacy/utility trade-off

For HBTS with $\alpha = 2^{-n}$ and $\beta = 0$, the confusion matrix is nearly identity—high utility but query results are observable (the “concession” discussed in Section III-C).

Throughout, we distinguish latent values (b, S, f) from observed approximations $(\tilde{b}, \tilde{S}, \tilde{f})$ carrying explicit error rates.

C. Threat Model

We consider an honest-but-curious adversary model where:

- Participants follow the protocol correctly but attempt to learn additional information
- The adversary has access to hash values but cannot invert the hash function
- The adversary may have auxiliary information about the data distribution
- The cryptographic hash function is modeled as a random oracle

We explicitly exclude:

- Malicious adversaries who deviate from the protocol
- Side-channel attacks on the implementation
- Quantum adversaries (though post-quantum hash functions could be used)

III. SYSTEM DESIGN

A. Core Abstractions

HBTS is built around three core abstractions that compose to enable complex privacy-preserving operations:

1) *Hash-Trapdoor Values*: A hash-trapdoor value is a one-way transformation of sensitive data: equality testing is preserved probabilistically (FPR equals hash collision probability, e.g., 2^{-256} for SHA-256) while original values cannot be recovered. Note that this bounds *collision errors*, not privacy—privacy depends on input entropy and resistance to dictionary attacks. Implementation details are in Appendix A.

2) *Approximate Values*: All operations in HBTS return approximate values carrying both the computed result and its false positive and false negative rates (α, β) , with confidence $1 - \alpha - \beta$ when $\alpha + \beta \leq 1$.

3) *Set Operations*: HBTS provides two primary set implementations with different algebraic properties:

Boolean Sets support full Boolean algebra:

- Union: $A \cup B$ with error propagation
- Intersection: $A \cap B$ with error composition
- Complement: $\neg A$ with error inversion
- Membership: $x \in A$ with hash collision probability

Symmetric Difference Sets form a group under XOR:

- XOR: $A \oplus B$ for disjoint unions
- Identity: Empty set
- Inverse: Every set is its own inverse
- Efficient for aggregation operations

B. Error Propagation

A key aspect of HBTS is systematic error propagation through operations. For Boolean operations, we derive tight bounds:

Theorem 3 (Union Error Bound). *For sets A and B with false positive rates α_A, α_B and false negative rates β_A, β_B :*

$$\alpha_{A \cup B} \leq \alpha_A + \alpha_B - \alpha_A \cdot \alpha_B \quad (1)$$

$$\beta_{A \cup B} = \beta_A \cdot \beta_B \quad (2)$$

Proof. Dual to Theorem 4 by De Morgan's laws: $A \cup B = \overline{A \cap B}$. Complement swaps error rates ($\alpha \leftrightarrow \beta$), and intersection multiplies false positive rates while accumulating false negative rates via inclusion-exclusion. Equivalently: a false positive in the union occurs if *either* set reports a false positive (independent events, yielding the inclusion-exclusion bound $\alpha_A + \alpha_B - \alpha_A \alpha_B$), while a false negative requires *both* sets to miss the element (yielding the product $\beta_A \cdot \beta_B$). $\square$

Theorem 4 (Intersection Error Bound). *For sets A and B with false positive rates α_A, α_B and false negative rates β_A, β_B :*

$$\alpha_{A \cap B} = \alpha_A \cdot \alpha_B \quad (3)$$

$$\beta_{A \cap B} \leq \beta_A + \beta_B - \beta_A \cdot \beta_B \quad (4)$$

Proof. For intersection, a false positive requires *both* sets to report false positives (independent events). A false negative

occurs if *either* set reports a false negative for an element truly in the intersection. $\square$

Note that intersection *reduces* false positive rates (multiplicative) while union *increases* them (sub-additive).

Theorem 5 (Boolean Error Composition). *For independent approximate Booleans $\tilde{b}_1, \tilde{b}_2$ with error rates (α_1, β_1) and (α_2, β_2) :*

$$\neg \tilde{b} : (\beta, \alpha) \text{ (rates swap)} \quad (5)$$

$$\tilde{b}_1 \wedge \tilde{b}_2 : (\alpha_1 \alpha_2, \beta_1 + \beta_2 - \beta_1 \beta_2) \quad (6)$$

$$\tilde{b}_1 \vee \tilde{b}_2 : (\alpha_1 + \alpha_2 - \alpha_1 \alpha_2, \beta_1 \beta_2) \quad (7)$$

Proof. Negation swaps the meaning of false positive and false negative. For conjunction, a false positive requires both operands to err (independent), while a false negative occurs if either errs. Disjunction is dual by De Morgan's laws. $\square$

Remark 6 (Bernoulli Model Foundation). *The closed-form error rates for AND, OR, NOT, and their compositions above are consequences of the Bernoulli model's element-wise independence axiom: each element's membership error is an independent channel. Formally, the joint confusion matrix for a t -element operation factors as a Kronecker (tensor) product of per-element confusion matrices $Q^{\otimes t}$, so that composed error rates reduce to scalar arithmetic on (α, β) pairs. This factorization is what makes the error propagation rules tractable—without independence, composed error bounds would require the full joint distribution over all elements.*

Corollary 7 (Composition Accumulation). *For t -fold composition of independent Bernoulli Booleans each with FPR α :*

- **t -fold OR (union)**: $\alpha_{\text{total}} = 1 - (1 - \alpha)^t \approx t\alpha$ for small α .
- **t -fold AND (intersection)**: $\alpha_{\text{total}} = \alpha^t$ (FPR decreases exponentially).

Dually, for FNR β : t -fold AND accumulates as $\beta_{\text{total}} = 1 - (1 - \beta)^t$, and t -fold OR reduces as $\beta_{\text{total}} = \beta^t$.

C. Layered Privacy: Trapdoor Representation, Bernoulli Results

HBTS exhibits a layered privacy model combining two distinct mechanisms: *trapdoor representation* through one-way hash transformations, and *Bernoulli query results* that are plaintext but approximate.

1) *Trapdoor Representation Layer*: The hash transformation $T_k(v) = H(k \| v)$ creates a *trapdoor representation*: inspecting a hash value reveals nothing about the underlying data beyond what can be learned through defined operations. This provides:

- **Preimage resistance**: Cannot recover v from $T_k(v)$ without exhaustive search
- **Key-dependent hiding**: Without knowledge of k , an adversary cannot verify membership guesses
- **Representation uniformity**: Hash values appear uniformly random, hiding structural properties of inputs

2) *Bernoulli Query Layer*: Operations on hash values return *Bernoulli Booleans*—plaintext approximate results with explicit error rates:

- **Membership queries**: $x \in S$ returns an observable true/false with rates $(\alpha, 0)$ where $\alpha = 2^{-n}$
- **Set operations**: Union and intersection return sets with composed error rates per Theorems 1–3
- **Results are not hidden**: An observer sees query outcomes as plaintext booleans

Remark 8 (The Concession). *While the trapdoor hides the underlying data, query results are observable. A fully access-pattern-hiding system would conceal even the pattern of true/false results (as in oblivious RAM [6]). HBTS trades this stronger property for practical efficiency—suitable when query results are already shared (e.g., private set intersection reveals intersection members) or when the pattern of results is not sensitive.*

3) *Hash Values as Bernoulli Set Carriers*: The `HashValue` type with bitwise operations ($\wedge, \vee, \oplus, \neg$) forms the carrier for Bernoulli set operations. Each hash equality test is a second-order Bernoulli Boolean:

$$\alpha = 2^{-n} \quad (\text{hash collision probability for } n\text{-bit hash}) \quad (8)$$

$$\beta = 0 \quad (\text{no false negatives for hash equality}) \quad (9)$$

This asymmetry ($\alpha > 0, \beta = 0$) propagates through composed operations via the error composition rules above.

D. Architecture

HBTS follows a layered architecture as shown in Figure 1:

Applications (PSI, Analytics, Aggregation)
Operations Layer (Similarity, Cardinality Estimation)
Set Layer (Boolean Algebra, Symmetric Difference)
Core Primitives (Hash-Trapdoor Values, Approximate Types)
Cryptographic Layer (SHA-256, BLAKE2b, etc.)

Fig. 1: HBTS layered architecture

IV. MATHEMATICAL FOUNDATIONS

A. Boolean Algebra Framework

HBTS is grounded in a homomorphism between two Boolean algebras: the *set algebra* $A := (\mathcal{P}(\{0, 1\}^*), \cup, \cap, \neg, \emptyset, \{0, 1\}^*)$ over arbitrary bit strings, and the *bit vector algebra* $B := (\{0, 1\}^m, \vee, \wedge, \neg, 0^m, 1^m)$ of m -bit vectors with bitwise operations.

Definition 9 (Trapdoor Homomorphism). *The homomorphism $F : A \rightarrow B$ is defined as:*

$$F(\beta) := \begin{cases} H(\beta) & \beta \in \{0, 1\}^* \\ \vee & \beta = \cup \\ \wedge & \beta = \cap \\ \neg & \beta = \neg_A \\ 0^m & \beta = \emptyset \\ 1^m & \beta = \{0, 1\}^* \end{cases}$$

where $H : \{0, 1\}^* \rightarrow \{0, 1\}^m$ is a cryptographic hash function. For a set $S = \{x_1, \dots, x_k\}$, this extends to $F(S) = H(x_1) \vee H(x_2) \vee \dots \vee H(x_k)$.

Caveat on intersection: The mapping $\cap \mapsto \wedge$ satisfies $F(A) \wedge F(B) \supseteq F(A \cap B)$, not equality. Extra 1-bits arise when an element of $A \setminus B$ and an element of $B \setminus A$ both set the same bit position, producing false positives in the intersection encoding. The union mapping $\cup \mapsto \vee$ is exact.

Remark 10 (Bloom Filter Connection). *The construction presented here is mathematically equivalent to a Bloom filter with a single hash function ($k = 1$). The contribution of this paper is not the construction itself, which is well known, but the algebraic analysis: the characterization as a Boolean algebra homomorphism, the precise error propagation formulas for composed operations, and the proof that complement is not preserved.*

Theorem 11 (One-Wayness of Trapdoor Homomorphism). *The homomorphism F is one-way for two independent reasons:*

- 1) **Non-injectivity**: F is not injective. Since $\{0, 1\}^*$ is countably infinite and $\{0, 1\}^m$ has only 2^m elements, by the pigeonhole principle, countably infinitely many inputs map to each output.
- 2) **Preimage resistance**: The hash function H is a cryptographic hash that resists preimage attacks. Given $y \in \{0, 1\}^m$, finding any x such that $H(x) = y$ requires $O(2^m)$ operations, while computing $H(x)$ from x is efficient.

Remark 12 (Bit-Rate Asymptotics). *The bit-rate (bits per element) required to achieve a target false positive rate ε for a set of s elements in an m -bit vector can be derived as follows. From the FPR formula $\varepsilon = \alpha(s)^m$ where $\alpha(s) = 1 - 2^{-(s+1)}$ is the per-bit occupancy probability, solving for m gives:*

$$m = \frac{\log_2 \varepsilon}{\log_2 \alpha(s)}$$

For large s , $\alpha(s) \rightarrow 1$ and writing $\delta = 2^{-(s+1)}$, we have:

$$\log_2(1 - \delta) \approx \frac{-\delta}{\ln 2}$$

so the required bit-vector length is:

$$m \approx \frac{-\log_2(\varepsilon) \cdot \ln 2}{\delta} = -\log_2(\varepsilon) \cdot \ln 2 \cdot 2^{s+1}$$

and the bit-rate m/s grows exponentially with s .

Remark 13 (Absolute Efficiency). *The absolute efficiency of the bit-vector representation is:*

$$\text{Eff}(s) = -s \log_2 \alpha(s)$$

which tends to 0 exponentially as $s \rightarrow \infty$ (since $\log_2 \alpha(s) \rightarrow 0$ while s grows linearly), making this construction practical only for small sets (typically $s \leq 20$). For larger sets, the two-level hashing scheme of Section IV-E is required.

B. Security Analysis

The privacy guarantee of HBTS is hash preimage resistance under the random oracle model.

Theorem 14 (Privacy Preservation). *Let $H : \{0, 1\}^* \rightarrow \{0, 1\}^n$ be modeled as a random oracle. Under the random oracle model, recovering x from $H(k||x)$ requires expected $O(2^n)$ hash evaluations when k is secret and uniformly random. No efficient algorithm can determine the input x from the hash output with non-negligible probability.*

Proof. In the random oracle model, $H(k||x)$ is uniformly distributed over $\{0, 1\}^n$ and independent of x from the adversary's perspective (since k is unknown). Each hash evaluation tests one candidate input, and the probability of matching the target output is 2^{-n} per evaluation. After q queries, the success probability is at most $q \cdot 2^{-n}$, which is negligible for q polynomial in n . $\square$

C. Approximate Algebraic Properties

HBTS exhibits approximate algebraic properties with explicit error bounds:

Definition 15 (Approximate Set Operations). *For hash-transformed sets $H(A)$ and $H(B)$ (represented as OR of per-element hash values), operations preserve set relationships as follows:*

- *Union:* $H(A) \vee H(B) = H(A \cup B)$ (exact by definition of bitwise OR)
- *Intersection:* $H(A) \wedge H(B) \supseteq H(A \cap B)$, i.e., the bitwise AND is a superset of the true intersection encoding. Equality holds when there are no cross-element hash collisions between $A \setminus B$ and $B \setminus A$. Each extra 1-bit arises with probability at most $|A \setminus B| \cdot |B \setminus A| / 2^n$ per bit position, yielding a per-element false positive rate bounded by $(1 - (1 - 2^{-n})^{|A \setminus B| \cdot |B \setminus A|})$.
- *Symmetric difference:* $H(A) \oplus H(B) = H(A \triangle B)$ (exact for disjoint sets)

D. Size-Dependent False Positive Analysis

The simple bound $\alpha = 2^{-n}$ assumes single-element queries. For sets of size k , the FPR depends on set density. We derive precise bounds from the bit-string representation.

Theorem 16 (Membership False Positive Rate). *For a set S of size k represented as n -bit strings, the false positive rate for membership queries is:*

$$\varepsilon_{\in}(k, n) = \left(1 - 2^{-(k+1)}\right)^n$$

For small k , this approximates to $\varepsilon_{\in} \approx e^{-n \cdot 2^{-(k+1)}}$.

Proof. After k insertions, each bit position is set to 1 with probability $1 - 2^{-k}$. For a random element $x \notin S$ to test positive, every bit position where $h(x)$ is 1 must also be 1 in $F(S)$. Since $h(x)$ has each bit set with probability $\frac{1}{2}$, the probability that bit j does not refute membership is:

$$\Pr[\neg(h(x)_j = 1 \wedge F(S)_j = 0)] = 1 - \frac{1}{2} \cdot 2^{-k} = 1 - 2^{-(k+1)}$$

The result follows from independence across n bit positions. $\square$

Theorem 17 (Subset False Positive Rate). *For sets S_1, S_2 of sizes k_1, k_2 , the false positive rate for the subset test $S_1 \subseteq S_2$ is:*

$$\varepsilon_{\subseteq}(k_1, k_2, n) = (1 - (1 - 2^{-k_1}) \cdot 2^{-k_2})^n$$

Proof. For $S_1 \subseteq S_2$ to test positive falsely, every bit set in $F(S_1)$ must be set in $F(S_2)$. Bit j is set in $F(S_1)$ with probability $1 - 2^{-k_1}$ and unset in $F(S_2)$ with probability 2^{-k_2} . The probability that bit j does not refute the subset relation is $1 - (1 - 2^{-k_1}) \cdot 2^{-k_2}$, and the result follows from independence. $\square$

Theorem 18 (Space Complexity). *To maintain a fixed false positive rate ε for membership queries on a set of size k , the hash size must satisfy:*

$$n = \mathcal{O}(2^k)$$

This exponential growth limits the single-level scheme to small sets (typically $k \leq 20$).

Proof. From Theorem 16, we have $\varepsilon = (1 - 2^{-(k+1)})^n$. Solving for n :

$$n = \frac{\log \varepsilon}{\log(1 - 2^{-(k+1)})} \approx \frac{\log \varepsilon}{-2^{-(k+1)}} = -\log(\varepsilon) \cdot 2^{k+1}$$

Thus $n = \mathcal{O}(2^k)$ for fixed ε . $\square$

Remark 19 (Information-Theoretic Consistency). *This exponential growth in k is consistent with the information-theoretic lower bound of $-\log_2(\varepsilon)$ bits per element from the Bernoulli entropy analysis. For the trapdoor Boolean algebra (online construction), the space budget n is fixed at construction time, forcing ε to degrade as k grows. Equivalently, each inserted element consumes $\mathcal{O}(1)$ bits of the fixed-size representation, leaving fewer bits to distinguish non-members, which is why ε increases monotonically with the set size k for fixed n .*

Theorem 20 (Complement Non-Preservation). *The homomorphism F does not preserve complement: $F(\neg_A x) \neq \neg_B F(x)$ for finite sets x .*

Proof. For any finite set x , the complement $\neg_A x = X^* \setminus x$ contains infinitely many elements from the free semigroup X^* . Since the hash function $h : X^* \rightarrow \{0, 1\}^n$ uniformly distributes over 2^n possible outputs, by the pigeonhole principle,




Fig. 2: False positive rate vs. hash size for different set sizes k . Left: small sets ($k = 4-10$) require bytes to achieve 5% FPR. Right: larger sets ($k = 12-16$) require kilobytes, demonstrating the $\mathcal{O}(2^k)$ space complexity.

$F(\neg_A x)$ will have all bit positions set to 1 as elements from $\neg_A x$ are added:

$$F(\neg_A x) = 0^n \mid h(y_1) \mid h(y_2) \mid \dots = 1^n$$

for any enumeration $\{y_1, y_2, \dots\} \subseteq \neg_A x$.

However, $\neg_B F(x) = \sim(h(x_1) \mid \dots \mid h(x_k))$ for finite $x = \{x_1, \dots, x_k\}$. Since $F(x)$ is a finite OR of k hash values, some bit positions will remain 0 in $F(x)$, making those positions 1 in $\neg_B F(x)$ but also making some positions 0. Thus $\neg_B F(x) \neq 1^n = F(\neg_A x)$. $\square$

Remark 21 (Structural Necessity of Complement Failure). *This asymmetry—AND/OR preserved (at least approximately), NOT not preserved—is structural, not an artifact of the hash construction. Any approximate Boolean algebra homomorphism from an infinite source algebra to a finite target algebra must fail to preserve complement, since the complement of a finite set in an infinite domain covers all hash values by pigeonhole. More precisely, for any finite codomain of size 2^m and any finite set x with $|\neg_A x| = \infty$, the image $F(\neg_A x)$ saturates to the maximal element 1^m , while $\neg_B F(x)$ is strictly submaximal whenever $F(x) \neq 0^m$. This obstruction is independent of the choice of hash function or encoding scheme.*

Remark 22 (Boolean Ring Alternative). *An alternative construction uses the Boolean ring $(\{0, 1\}^m, \oplus, \wedge, \text{id}, 0^m)$ with symmetric difference (XOR) instead of union. This construction maps sets to hash values via $G(\{x_1, \dots, x_k\}) = h(x_1) \oplus \dots \oplus h(x_k)$. The ring structure supports only equality predicates—membership, subset, and complement operations are not available. However, equality testing has optimal FPR of 2^{-m} independent of set size, and the output distribution is perfectly uniform, providing stronger privacy guarantees for equality-only applications.*

E. Two-Level Hashing for Scalability

The exponential space complexity of single-level hashing (Theorem 18) motivates a hierarchical approach. The two-level scheme partitions elements into bins, reducing effective set size per bin.

Definition 23 (Two-Level Hash Structure). *Given hash output size q bits, partition into:*




Fig. 3: Two-level hashing achieves practical FPR for large sets ($k = 15,000$) with approximately 60KB storage.

- *First w bits: bin index (selecting one of 2^w bins)*
 - *Remaining $q - w$ bits: element representation within bin*
- Total space: $n = 2^w(q - w)$ bits. Expected elements per bin: $k/2^w$.

Theorem 24 (Two-Level FPR). *The membership FPR for the two-level scheme with k elements is:*

$$\varepsilon(k, w, q) = \left(1 - 2^{-(k/2^{w+1})}\right)^{q-w}$$

This achieves practical FPR for large sets: with $w = 8$, $q = 256$, and $k = 1000$, we have approximately 4 elements per bin and $\text{FPR} \approx 2^{-248}$, using only 63.5 KB total storage.

F. Unified Hash Construction

All HBTS operations derive from a general hash-based construction that unifies probabilistic data structures with cryptographic privacy [4], [7]:

Definition 25 (Hash-Based Bernoulli Type). *A hash-based Bernoulli type $\mathcal{B}\langle T \rangle$ over base type T consists of:*

- 1) *A keyed hash function $h : \{0, 1\}^* \times \{0, 1\}^s \rightarrow \{0, 1\}^m$*
- 2) *For each output value y , a set of valid encodings $\text{Valid}(y) \subseteq \{0, 1\}^m$*
- 3) *A seed s such that for all inputs x : $h(\text{enc}(x), s) \in \text{Valid}(f(x))$*

The resulting type is a Bernoulli approximation of T : operations on $\mathcal{B}\langle T \rangle$ approximate operations on T with explicit error rates.

The false positive rate emerges from encoding set sizes: $\alpha = |\text{Valid}(\text{true})|/2^m$. This unified view reveals that Bloom filters, trapdoor functions, and approximate maps are all instances of the same construction with different encoding strategies:

- **Bloom filter:** Multiple hash functions create overlapping valid encoding regions, yielding $\alpha = (1 - e^{-kn/m})^k$
- **HBTS trapdoor:** Single keyed hash with $\text{Valid}(\text{true}) = \{H(k\|v)\}$, yielding $\alpha = 2^{-m}$ (collision probability)
- **Approximate maps:** Encoding sizes vary by output value, enabling homophonic (frequency-proportional) allocation for privacy

G. Privacy-Space Trade-off

The same hash construction achieves fundamentally different goals by varying encoding sizes [4]:

Space-Optimal Encoding: $|\text{Valid}(y)| = 1$ (one encoding per value). Minimizes the total number of valid cipher values, but the single representation of each y is then observed with that value's own frequency, revealing the frequency distribution. Used in compression and space-efficient data structures.

Privacy-Optimal Encoding: $|\text{Valid}(y)| \propto \text{freq}(y)$ (proportional to frequency). Frequent values receive proportionally more encodings, so under the real query distribution each individual cipher value is equiprobable and the observed cipher-value distribution is flat, hiding frequency patterns. This is the classical Simmons homophonic prescription [8], the key insight connecting approximation to privacy.

Theorem 26 (Uniformity from Homophonic Allocation). *Let each output value y occur with probability $\text{freq}(y)$ and be encoded by a uniform choice within $\text{Valid}(y)$. With encoding-set sizes $|\text{Valid}(y)| = c \cdot \text{freq}(y)$ for constant c , every cipher value $v \in \text{Valid}(y)$ is observed with probability*

$$P[\text{cipher value} = v] = \frac{\text{freq}(y)}{|\text{Valid}(y)|} = \frac{1}{N}, \quad N := \sum_y |\text{Valid}(y)|$$

so the observed cipher-value distribution is uniform without explicit randomization. (The inverse allocation $|\text{Valid}(y)| \propto 1/\text{freq}(y)$ sharpens the distribution instead, the opposite of the privacy goal.)

HBTS uses uniform-sized encodings (a middle ground), providing practical privacy at constant space cost per element. Each element receives a fixed-size hash representation, offering computational privacy when the input domain has sufficient entropy.

H. Cardinality Estimation

HBTS incorporates the well-established HyperLogLog algorithm [9] for cardinality estimation on hash-transformed sets. We apply HyperLogLog without modification, relying on its proven accuracy guarantees:

Algorithm 1 Cardinality Estimation

Require: Trapdoor set S

Ensure: Estimated cardinality $\hat{n}$

```

1:  $m \leftarrow$  number of buckets
2:  $M \leftarrow$  array of  $m$  registers
3: for each trapdoor  $t \in S$  do
4:    $j \leftarrow$  first  $\log_2 m$  bits of  $t$ 
5:    $w \leftarrow$  remaining bits of  $t$ 
6:    $M[j] \leftarrow \max(M[j], \rho(w))$ 
7: end for
8:  $\hat{n} \leftarrow \alpha_m \cdot m^2 / \sum_{j=1}^m 2^{-M[j]}$ 
9: return  $\hat{n}$ 
```

The algorithm achieves relative error $1.04/\sqrt{m}$ using $O(m \log \log n)$ bits.

V. IMPLEMENTATION

We implemented HBTS as a Python library (`cts`, approximately 200 lines of code) using only the standard library. The implementation provides four core modules:

- **HashValue:** Fixed-size byte arrays with bitwise Boolean operations (AND, OR, XOR, NOT), providing the carrier for Bernoulli set operations.
- **BernoulliBoolean:** Boolean values carrying explicit error rates (α, β) , with confusion matrix computation, Bayesian posterior updates, and mutual information calculation.
- **Trapdoor/TrapdoorFactory:** One-way transformation $T_k(v) = H(k||v)$ using HMAC-SHA256, with key fingerprinting for compatibility checking across sets.
- **BooleanSet:** Full Boolean algebra over sets of trapdoors, supporting union ($\cup$), intersection ($\cap$), symmetric difference ($\hat{\cup}$), and set difference ($-$), with automatic error rate propagation and cardinality bound tracking through operations.

The library requires Python 3.10+ and has no external dependencies. All operations return **BernoulliBoolean** results with explicit error rates, and set operations propagate error bounds using the composition rules from Section III. Implementation details including code listings are provided in Appendix A.

VI. EVALUATION

A. Complexity Analysis

The asymptotic complexity of HBTS operations is determined by the underlying hash computation and set operations:

TABLE I: Operation Complexity

Operation	Time	Space
Trapdoor creation	$O(1)$	$O(n)$ bits
Set membership test	$O(1)$ expected	$O(1)$
Set intersection (k_1, k_2)	$O(\min(k_1, k_2))$	$O(\min(k_1, k_2))$
Set union (k_1, k_2)	$O(k_1 + k_2)$	$O(k_1 + k_2)$

All operations are dominated by hash computation cost. Since HBTS uses standard cryptographic hashing (SHA-256), individual hash evaluations complete in the sub-microsecond range on modern hardware, making HBTS orders of magnitude faster than homomorphic encryption or MPC-based alternatives for set operations.

B. Security Evaluation

The security of HBTS rests on the collision resistance and preimage resistance of the underlying hash function. For SHA-256, the expected collision probability is 2^{-256} per pair, providing a negligible false positive rate for practical set sizes.

C. Limitations and Attack Vectors

HBTS provides practical privacy but has important limitations that users must understand:

Dictionary Attacks: For low-entropy input domains (e.g., phone numbers, SSNs, email addresses), an adversary with

knowledge of the domain can precompute hashes for all possible values and match against observed trapdoors. Privacy degrades to $O(1/|D|)$ where $|D|$ is the domain size. *Mitigation*: Use high-entropy inputs, apply key stretching (e.g., Argon2), or combine with additional randomization.

Frequency Leakage: Query frequencies are preserved through hash transformations. An adversary observing repeated queries learns which trapdoors correspond to frequently-accessed values. *Mitigation*: Add dummy queries, use query batching, or apply differential privacy noise.

Correlation Leakage: Deterministic hashing reveals equality patterns—repeated queries on the same value produce identical hashes. This leaks that certain queries target the same underlying value. *Mitigation*: Periodic key rotation, though this complicates system design.

Quantum Vulnerability: Grover’s algorithm reduces effective hash security from n bits to $n/2$ bits. For quantum resistance, use 512-bit hashes to maintain 256-bit security.

Remark 27 (When HBTS is Appropriate). *HBTS is suitable when: (1) input domains have high entropy, (2) dictionary attacks are infeasible, (3) frequency patterns are acceptable leakage, or (4) approximate plausible deniability suffices. For stronger guarantees, consider FHE or MPC despite their performance costs.*

D. Comparison with Alternatives

TABLE II: Comparison with Related Systems

System	Privacy Model	Dictionary Resistant	Performance
HBTS	Preimage	No*	μ s
FHE	Semantic	Yes	s
MPC	Simulation	Yes	ms
Bloom Filters	None	N/A	μ s

*Vulnerable when input domain has low entropy

HBTS occupies a practical niche: it provides preimage-based privacy with microsecond performance, suitable for high-entropy domains where dictionary attacks are infeasible.

VII. APPLICATIONS

A. Private Set Intersection

HBTS enables efficient PSI without revealing non-matching elements. Using hash-table based intersection, HBTS achieves $O(n)$ complexity with $O(n)$ space, matching the best non-private approaches. The key advantage is that operations occur entirely in the hash domain—the false positive rate for membership is bounded by hash collision probability, while the underlying values remain computationally hidden.

B. Secure Deduplication

Cloud storage providers can identify duplicate files without accessing content:

Algorithm 2 Secure Deduplication

Require: File F , Trapdoor factory T

- 1: $chunks \leftarrow$ split F into blocks
 - 2: $hashes \leftarrow$ empty set
 - 3: **for** each chunk $c \in chunks$ **do**
 - 4: $h \leftarrow T.create(c)$
 - 5: add h to $hashes$
 - 6: **end for**
 - 7: Query storage for existing $hashes$
 - 8: Upload only unique chunks
-

C. Privacy-Preserving Analytics

HBTS supports various analytical operations on hash-transformed data:

Histogram Generation: Count occurrences without revealing underlying values, useful for distribution analysis while preserving privacy.

Frequency Analysis: Identify common patterns in hash-transformed data with collision-bounded error rates.

Similarity Metrics: Compute Jaccard similarity and other set-based metrics on private sets with explicit error quantification.

D. Federated Learning

HBTS supports secure aggregation in federated learning by allowing participants to submit hash-transformed model updates. The server aggregates these trapdoor-preserving updates using symmetric difference operations, revealing only the final aggregate while preserving individual update privacy.

This approach is particularly effective when combined with differential privacy for additional statistical guarantees.

VIII. RELATED WORK

A. Homomorphic Encryption

Fully homomorphic encryption (FHE) enables arbitrary computation on encrypted data. Since Gentry’s breakthrough [1], significant progress has been made in practical FHE systems. TFHE [10] and CKKS [11] achieve sub-second bootstrapping, while Brakerski’s FHE without modulus switching [12] simplifies implementation. However, FHE still incurs 1000-10000 $\times$ overhead for general computation. Partially homomorphic schemes like Paillier [13] offer better performance but support only specific operations. HBTS provides a complementary approach, achieving microsecond performance by accepting approximate results rather than exact homomorphic computation.

B. Secure Multi-Party Computation

MPC protocols enable joint computation without revealing inputs. Classical protocols [2], [14] laid theoretical foundations, while modern frameworks have made MPC practical. MP-SPDZ [15] provides a comprehensive toolkit supporting multiple protocols, while ABY3 [16] achieves efficient three-party computation. Recent advances in silent OT [17] and function secret sharing [18] have reduced communication

complexity. However, MPC still requires multiple rounds of interaction and coordination between parties. HBTS operates non-interactively using only hash transformations, trading exact computation for practical single-party performance.

C. Private Set Intersection

PSI protocols have evolved significantly since early work [19], [20]. Modern protocols achieve remarkable efficiency: OPRF-based PSI [21] handles billion-element sets, while circuit-PSI [22] enables arbitrary computations on intersection results. Unbalanced PSI [23] optimizes for asymmetric set sizes common in practice. Recent work on PSI from pseudorandom correlation generators [24] achieves optimal communication. HBTS provides a simpler alternative where approximate results are acceptable, requiring only hash computation without cryptographic protocols.

D. Approximate Data Structures

Probabilistic data structures trade accuracy for efficiency. Bloom filters [5] pioneered this approach for membership testing. Cuckoo filters [25] improve on Bloom filters by supporting deletions. Count-Min sketches [26] and streamed cardinality estimation [27] enable frequency and cardinality estimation. Recent work includes learned Bloom filters [28] using machine learning to optimize performance, and XOR filters [29] achieving near-optimal space efficiency. HBTS builds upon these foundations, adding cryptographic privacy through hash transformations while maintaining the efficiency benefits of approximate operations.

E. Differential Privacy

Differential privacy (DP) [3] provides statistical privacy guarantees through calibrated noise addition. The field has matured significantly with deployment at major tech companies [30], [31]. Recent advances include mechanism design via differential privacy [32], concentrated DP [33] for tighter privacy accounting, and shuffle DP [34] amplifying privacy through anonymization. Practical secure aggregation [35] enables federated learning at scale. HBTS provides complementary cryptographic privacy that can be composed with DP techniques, offering defense-in-depth for sensitive applications.

IX. CONCLUSION

We presented Hash-Based Trapdoor Sets (HBTS), a practical framework for privacy-preserving set operations that combines cryptographic hash functions with probabilistic data structures. By explicitly managing error rates and embracing approximation, HBTS achieves efficient performance—bounded by the cost of standard hash computations—while providing privacy bounded by hash collision probabilities.

Our contributions include:

- A systematic framework for error propagation through composed set operations on hash-transformed data
- Integration of established probabilistic algorithms (Bloom filters, HyperLogLog) with cryptographic privacy guarantees

- A reference Python implementation with operations reducible to standard hash computations
- Validation through real-world applications in private set intersection, secure aggregation, and federated learning

Future directions include:

- Integration with post-quantum hash functions for quantum resistance
- Hardware acceleration via AES-NI and SHA extensions
- Composition with differential privacy for enhanced protection
- Formal verification of implementation correctness

A. Current Limitations

The current Python reference implementation provides core trapdoor and set operations with full error propagation. The following features are not yet implemented:

- **Cryptographic key derivation:** The implementation uses direct SHA-256 concatenation ($H(k||v)$); production deployment would benefit from HKDF or Argon2 for proper key derivation
- **Differential privacy integration:** Combining HBTS with calibrated noise addition is a natural extension but not yet implemented
- **Constant-time primitives:** The Python implementation does not guarantee constant-time execution; side-channel resistance would require a lower-level language
- **Performance optimization:** The reference implementation prioritizes clarity over performance; a production implementation in C or Rust would benefit from SIMD hashing and memory pooling

HBTS demonstrates that practical privacy-preserving computation is achievable when applications can accept approximate results with explicit error bounds.

REFERENCES

- [1] C. Gentry, “Fully homomorphic encryption using ideal lattices,” in *Proceedings of the 41st annual ACM symposium on Theory of computing*, 2009, pp. 169–178.
- [2] A. C.-C. Yao, “Protocols for secure computations,” in *23rd Annual Symposium on Foundations of Computer Science*. IEEE, 1982, pp. 160–164.
- [3] C. Dwork, F. McSherry, K. Nissim, and A. Smith, “Calibrating noise to sensitivity in private data analysis,” in *Theory of Cryptography Conference*. Springer, 2006, pp. 265–284.
- [4] A. Towell, “Bernoulli types: A framework for approximate computation,” 2024, working paper.
- [5] B. H. Bloom, “Space/time trade-offs in hash coding with allowable errors,” *Communications of the ACM*, vol. 13, no. 7, pp. 422–426, 1970.
- [6] X. S. Wang, K. Nayak, C. Liu, H. Chan, E. Shi, E. Stefanov, and Y. Huang, “Oblivious data structures,” *Proceedings of the ACM SIGSAC Conference on Computer and Communications Security*, pp. 215–226, 2014.
- [7] A. Towell, “Trapdoor computing: Privacy through uniform encoding,” 2024, working paper.
- [8] G. J. Simmons, “Symmetric and asymmetric encryption,” *ACM Computing Surveys*, vol. 11, no. 4, pp. 305–330, 1979.
- [9] P. Flajolet, E. Fusy, O. Gandouet, and F. Meunier, “Hyperloglog: the analysis of a near-optimal cardinality estimation algorithm,” in *Conference on Analysis of Algorithms*, 2007, pp. 127–146.
- [10] I. Chillotti, N. Gama, M. Georgieva, and M. Izabachène, “Tfhe: Fast fully homomorphic encryption over the torus,” vol. 33, no. 1, 2020, pp. 34–91.

- [11] J. H. Cheon, A. Kim, M. Kim, and Y. Song, “Homomorphic encryption for arithmetic of approximate numbers,” in *International Conference on the Theory and Application of Cryptology and Information Security*, 2017, pp. 409–437.
- [12] Z. Brakerski, “Fully homomorphic encryption without modulus switching from classical GapSVP,” in *Annual Cryptology Conference (CRYPTO)*. Springer, 2012, pp. 868–886.
- [13] P. Paillier, “Public-key cryptosystems based on composite degree residuosity classes,” in *International Conference on the Theory and Application of Cryptographic Techniques*. Springer, 1999, pp. 223–238.
- [14] O. Goldreich, S. Micali, and A. Wigderson, “How to play any mental game,” in *Proceedings of the 19th Annual ACM Symposium on Theory of Computing*, 1987, pp. 218–229.
- [15] M. Keller, “Mp-spdz: A versatile framework for multi-party computation,” in *Proceedings of the 2020 ACM SIGSAC Conference on Computer and Communications Security*, 2020, pp. 1575–1590.
- [16] P. Mohassel and P. Rindal, “Aby3: A mixed protocol framework for machine learning,” in *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security*, 2018, pp. 35–52.
- [17] E. Boyle, G. Couteau, N. Gilboa, and Y. Ishai, “Cheaper silent ot extension and applications to multi-server pir,” *Annual International Cryptology Conference*, pp. 371–403, 2022.
- [18] E. Boyle, N. Chandran, N. Gilboa, D. Gupta, Y. Ishai, N. Kumar, and M. Rathee, “Function secret sharing for mixed-mode and fixed-point secure computation,” in *Annual International Conference on the Theory and Applications of Cryptographic Techniques*, 2021, pp. 871–900.
- [19] C. Meadows, “A more efficient cryptographic matchmaking protocol for use in the absence of a continuously available third party,” *Proceedings of IEEE Symposium on Security and Privacy*, pp. 134–137, 1986.
- [20] M. J. Freedman, K. Nissim, and B. Pinkas, “Efficient private matchings and set intersection,” in *International Conference on the Theory and Applications of Cryptographic Techniques*. Springer, 2004, pp. 1–19.
- [21] S. Raghuraman and P. Rindal, “Blazing fast psi from improved okvs and subfield vole,” in *Proceedings of the 2022 ACM SIGSAC Conference on Computer and Communications Security*, 2022, pp. 2505–2517.
- [22] N. Chandran, D. Gupta, and A. Shah, “Circuit-psi with linear complexity via relaxed batch oprf,” in *Proceedings on Privacy Enhancing Technologies*, vol. 2022, no. 1, 2022, pp. 353–372.
- [23] H. Chen, K. Laine, and P. Rindal, “Fast private set intersection from homomorphic encryption,” in *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*, 2017, pp. 1243–1255.
- [24] P. Schoppmann, A. Gascon, and F. Kerschbaum, “Psi from pseudorandom correlation generators,” *Annual International Cryptology Conference*, pp. 332–364, 2023.
- [25] B. Fan, D. G. Andersen, M. Kaminsky, and M. D. Mitzenmacher, “Cuckoo filter: Practically better than bloom,” in *Proceedings of the 10th ACM International Conference on emerging Networking Experiments and Technologies*, 2014, pp. 75–88.
- [26] G. Cormode and S. Muthukrishnan, “An improved data stream summary: the count-min sketch and its applications,” *Journal of Algorithms*, vol. 55, no. 1, pp. 58–75, 2005.
- [27] D. Ting, “Streamed approximate counting of distinct elements: Beating optimal batch methods,” in *Proceedings of the 20th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2014, pp. 442–451.
- [28] T. Kraska, A. Beutel, E. H. Chi, J. Dean, and N. Polyzotis, “The case for learned index structures,” in *Proceedings of the 2018 International Conference on Management of Data*, 2018, pp. 489–504.
- [29] T. M. Graf and D. Lemire, “Xor filters: Faster and smaller than bloom and cuckoo filters,” vol. 25, 2020, pp. 1–16.
- [30] Ú. Erlingsson, V. Pihur, and A. Korolova, “Rappor: Randomized aggregatable privacy-preserving ordinal response,” in *Proceedings of the 2014 ACM SIGSAC Conference on Computer and Communications Security*, 2014, pp. 1054–1067.
- [31] A. D. P. Team, “Learning with privacy at scale,” *Apple Machine Learning Journal*, vol. 1, no. 8, 2017.
- [32] F. McSherry and K. Talwar, “Mechanism design via differential privacy,” in *48th Annual IEEE Symposium on Foundations of Computer Science (FOCS)*. IEEE, 2007, pp. 94–103.
- [33] M. Bun and T. Steinke, “Concentrated differential privacy: Simplifications, extensions, and lower bounds,” in *Theory of Cryptography Conference*, 2016, pp. 635–658.
- [34] Ú. Erlingsson, V. Feldman, I. Mironov, A. Raghunathan, K. Talwar, and A. Thakurta, “Amplification by shuffling: From local to central differential privacy via anonymity,” in *Proceedings of the Thirtieth Annual ACM-SIAM Symposium on Discrete Algorithms*, 2019, pp. 2468–2479.
- [35] K. Bonawitz, V. Ivanov, B. Kreuter, A. Marcedone, H. B. McMahan, S. Patel, D. Ramage, A. Segal, and K. Seth, “Practical secure aggregation for privacy-preserving machine learning,” in *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*, 2017, pp. 1175–1191.

APPENDIX

This appendix provides implementation details for the `cts` Python library that accompanies this paper. The library is approximately 200 lines of code with no external dependencies beyond the Python standard library.

A. Core Data Structures

1) *Trapdoor Implementation*: The trapdoor factory creates one-way transformations using keyed SHA-256:

Listing 1: Trapdoor factory (trapdoor.py)

```
class TrapdoorFactory:
    def __init__(self, key=None, hash_size=32):
        self._key = key or secrets.token_bytes(32)
        self._hash_size = hash_size
        self._fingerprint = hashlib.sha256(
            self._key).digest()[8:]

    def create(self, value):
        """T_k(v) = H(k || v)"""
        if isinstance(value, str):
            value = value.encode('utf-8')
        h = hashlib.sha256(self._key + value)
        return Trapdoor(
            hash=HashValue(h.digest()[self._hash_size:]),
            key_fingerprint=self._fingerprint)
```

Listing 2: Bernoulli Boolean with error tracking (approximate.py)

2) *Bernoulli Boolean Implementation*:

```
@dataclass(frozen=True, slots=True)
class BernoulliBoolean:
    """Boolean with explicit error rates."""
    value: bool
    alpha: float = 0.0 # FPR
    beta: float = 0.0 # FNR

    @property
    def confidence(self):
        return max(0.0, 1.0 - self.alpha - self.beta)

    @property
    def confusion_matrix(self):
        return ((1 - self.alpha, self.alpha),
                (self.beta, 1 - self.beta))
```

B. Error Composition Functions

Listing 3: Error propagation rules (approximate.py)

```
def compose_and(a1, b1, a2, b2):
    """AND: FPR = product, FNR = inclusion-exclusion"""
    return (a1 * a2, b1 + b2 - b1 * b2)
```

```

5 def compose_or(a1, b1, a2, b2):
6     """OR: FPR = inclusion-exclusion, FNR = product.
       """
7     return (a1 + a2 - a1 * a2, b1 * b2)
8
9 def compose_not(alpha, beta):
10    """NOT: rates swap."""
11    return (beta, alpha)

```

C. Boolean Set Operations

Listing 4: Set operations with error propagation (boolean_set.py)

```

1 class BooleanSet:
2     def __or__(self, other):
3         """Union with FPR propagation."""
4         self._check_compatible(other)
5         new_alpha, new_beta = compose_or(
6             self.alpha, self.beta,
7             other.alpha, other.beta)
8         return BooleanSet(
9             _elements=self._elements | other.
10                _elements,
11             alpha=new_alpha, beta=new_beta,
12             _low_k=max(self._low_k, other._low_k),
13             _high_k=self._high_k + other._high_k)
14
15     def __and__(self, other):
16         """Intersection with FPR propagation."""
17         self._check_compatible(other)
18         new_alpha, new_beta = compose_and(
19             self.alpha, self.beta,
20             other.alpha, other.beta)
21         return BooleanSet(
22             _elements=self._elements & other.
23                _elements,
24             alpha=new_alpha, beta=new_beta,
25             _low_k=0,
26             _high_k=min(self._high_k, other._high_k)
27         )

```